

# Applied Data Analysis (CS401)

Lecture 13

Scaling to  
data

13 Dec 2023



**model**  
**massive**

**EPFL**

**Robert West**



# Announcements

- Homework H2
  - Feedback to be released later this week
  - Reminder (see [Ed](#)): Please participate in [ML4Ed study](#) by Fri 14:00
- Project milestone P3 due next week (Fri 22 Dec)
- Friday's lab session:
  - Last lab session! → Last quiz (on lecture 12)
  - Project office hour (same [sign-up protocol](#) as last week)
  - Exercises on Spark (useful for your future projects, your job, your love life)
- **Course eval is available on IS-Academia!**
  - Note: different from the eval from a few weeks ago!

# Feedback

Give us feedback on this lecture here:

<https://go.epfl.ch/ada2023-lec13-feedback>

- What did you (not) like about this lecture?
- What was (not) well explained?
- On what would you like more (fewer) details?
- Where is Waldo?
- ...



## So far in this class...

- We made one big assumption:
  - All data fits on a single machine
  - Even more, all data fits into memory on a single machine (Pandas)
- Realistic assumption for **prototyping**, but frequently not for production code

# The big-data problem

Data is growing faster than computation speed

- + Growing data sources  
(e.g, Web, mobile, sensors, ...)
- + Cheap hard-disk storage
- Stalling CPU speeds
- RAM bottlenecks



# Examples

Facebook's daily logs: 60 TB

1000 Genomes project: 200 TB

Google Web index: [100+ PB](#)

Cost of 1 TB of disk: \$50

Time to read 1 TB from disk: 3 hours (100 MB/s)



**DISCLAIMER**

**These numbers  
(anno domini  
2016) are  
outdated (too  
small)!**

# The big-data problem

A single machine can no longer store, let alone process, all the data

The only solution is to **distribute** over a large cluster of machines

# But how much data should you get?

Of course, “it depends”, but for many applications the answer is:

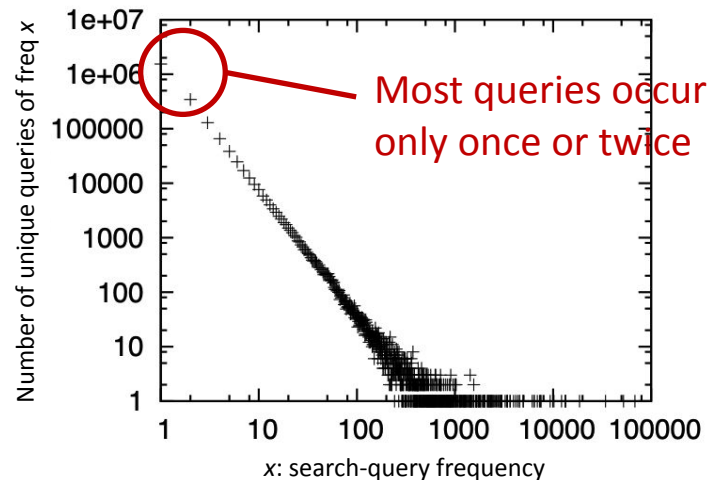
**As much as you can get**

Big data about people (text, Web, social media) tends to follow heavy-tailed distributions

(e.g., power laws)

Example: Web search

59% of all Web search queries are unique  
17% of all queries are made only twice  
8% are made three times





# Hardware for big data

**Budget** (a.k.a. commodity) hardware  
Not “gold-plated” (a.k.a. custom)


Many low-end servers

**Easy to add capacity**

**Cheaper** per CPU and per disk

**Increased complexity in software:**

- Fault tolerance
- Virtualization (e.g., distributed file systems)



# Problems with cheap hardware

**Failures**, e.g. (Google numbers)

- 1–5% hard drives/year
- 0.2% DIMMs (dual in-line memory modules)/year

**Commodity network** (1–10 Gb/s) speeds vs. RAM

- Much more latency (100–100,000x)
- Lower throughput (100–1,000x)

**Uneven performance**

- Inconsistent hardware (e.g., old + new)
- Variable network latency
- External loads



**DISCLAIMER**

These numbers are constantly changing thanks to new technology!



# Google datacenter

The background image shows a vast, industrial-scale data center. Rows of server racks are visible, illuminated by blue and yellow lights. The ceiling is high with a complex network of steel beams and pipes. The overall atmosphere is one of a high-tech, secure environment.


How to program this thing?

What's hard about cluster computing?

- 1. How to split work across machines?**
- 2. How to deal with failures?**

# How do you count the number of occurrences of each word in a document?


"I am Sam  
I am Sam  
Sam I am  
Do you like  
Green eggs and  
ham?"



I: 3  
am: 3  
Sam: 3  
do: 1  
you: 1  
like: 1  
...

# A hashtable (a.k.a. dict)!

"I am Sam  
I am Sam  
Sam I am  
Do you like  
Green eggs and  
ham?"



# A hashtable!

"I am Sam

I am Sam

Sam I am

Do you like

Green eggs and  
ham?"

{I: 1}

# A hashtable!

"I am Sam

I am Sam

Sam I am

Do you like

Green eggs and  
ham?"

{I: 1,  
am: 1}



# A hashtable!

"I am Sam  
I am Sam  
Sam I am  
Do you like  
Green eggs and  
ham?"

{I: 1,  
am: 1,  
Sam: 1}

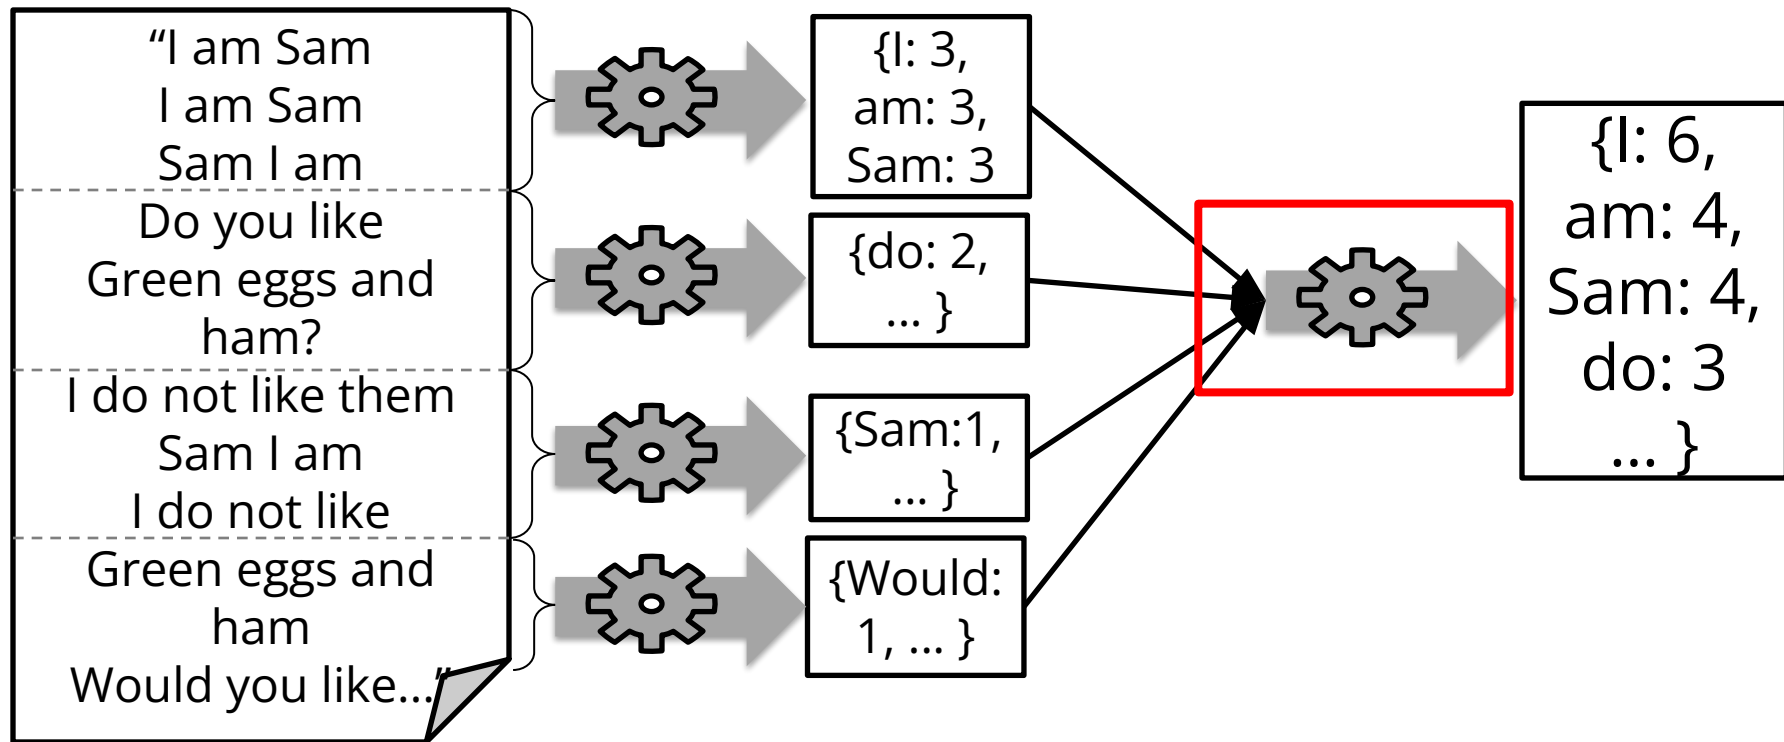
# A hashtable!

"I am Sam  
I am Sam  
Sam I am  
Do you like  
Green eggs and  
ham?"

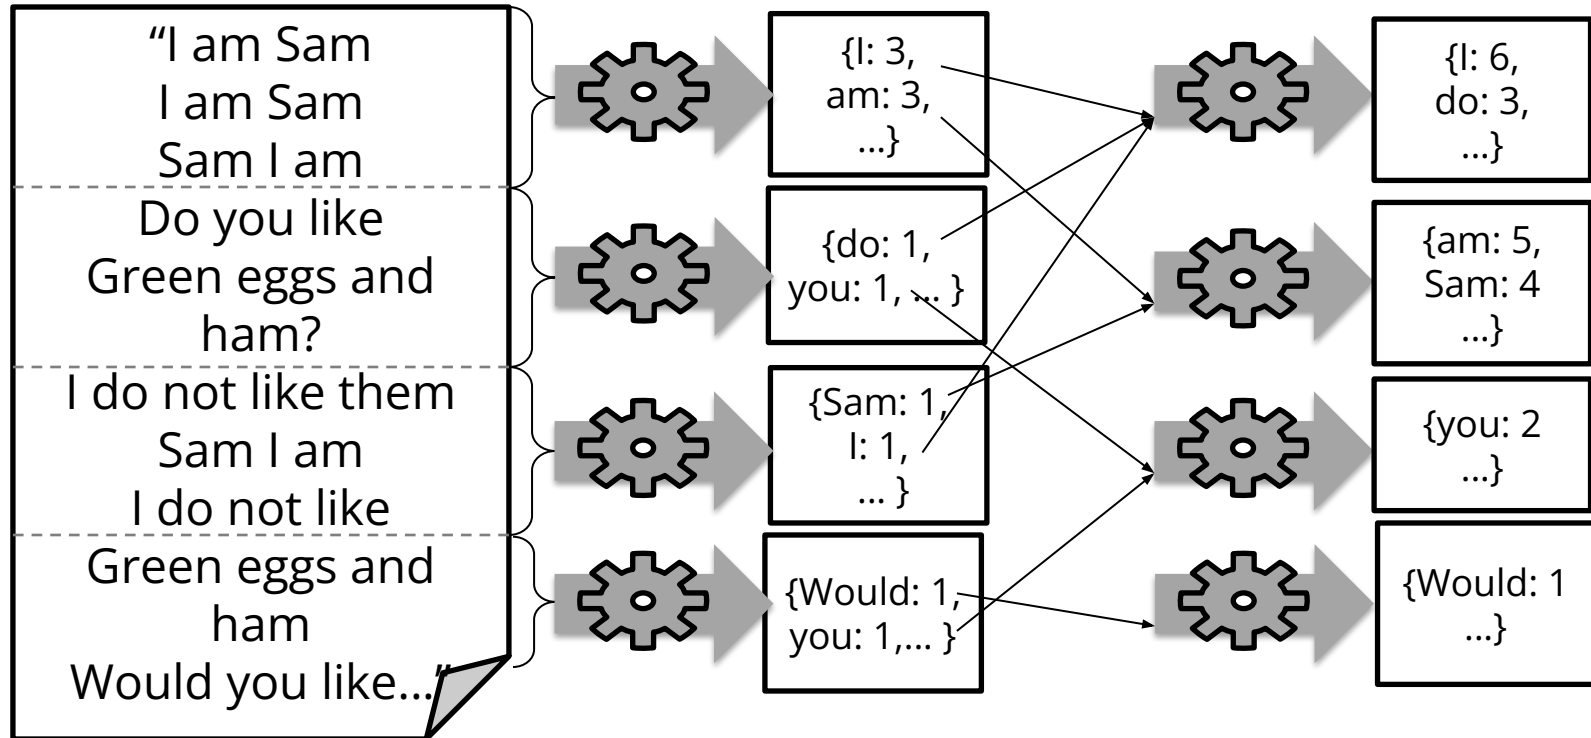
{I: 2,  
am: 1,  
Sam: 1}

What if the document is really  
big?

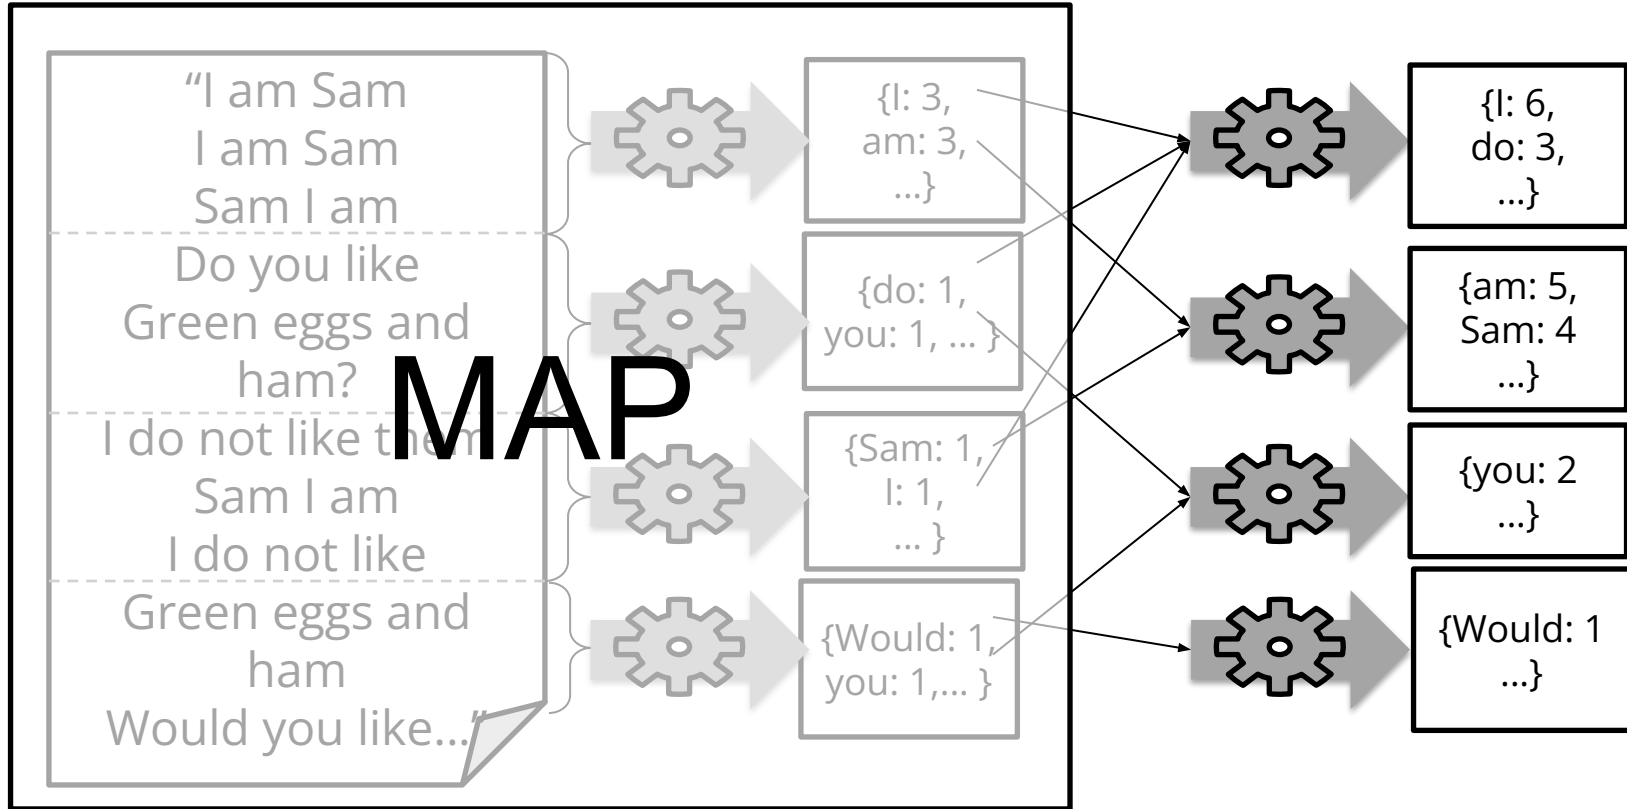
# What if the document is really big?




# "Divide and Conquer"



# “Divide and Conquer”



# “Divide and Conquer”



# Recall: What's hard about cluster computing?

## 1. How to divide work across machines?

- Moving data may be very expensive
- Must consider network, data locality

## 2. How to deal with failures?

- 1 server fails every 3 years  $\Rightarrow$  10K servers see ~10 faults/day
- Even worse: stragglers (node not failed, but slow)




# Solution: MapReduce

- Smart systems engineers have done all the work for you
  - Task scheduling
  - Virtualization of file system
  - Fault tolerance (incl. data replication)
  - Job monitoring
  - etc.
- “All” you need to do: implement Mapper and Reducer classes



Jeff Dean [\[facts\]](#)

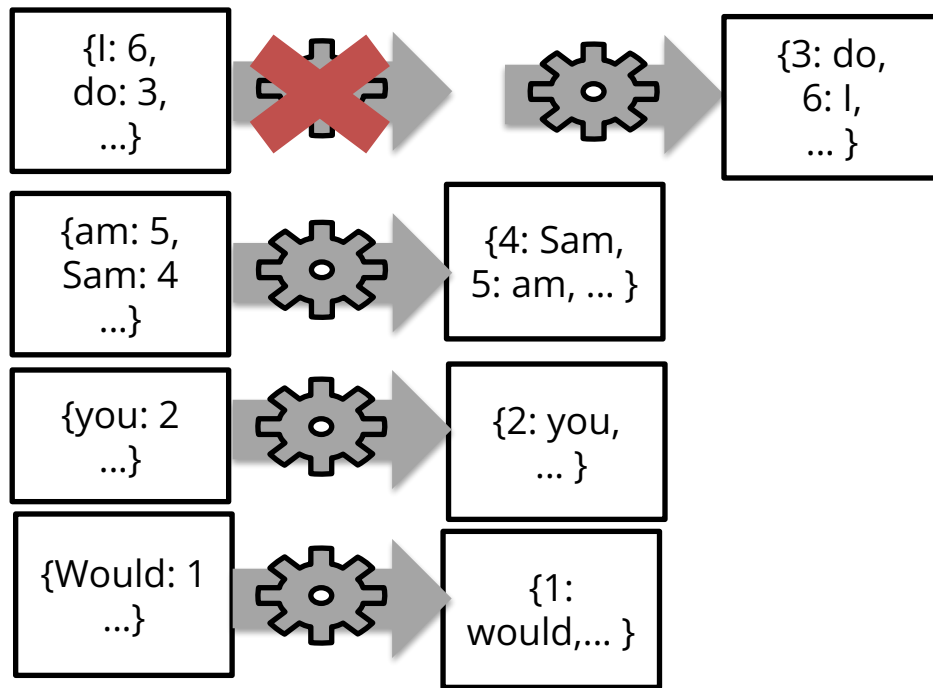
**hadoop**





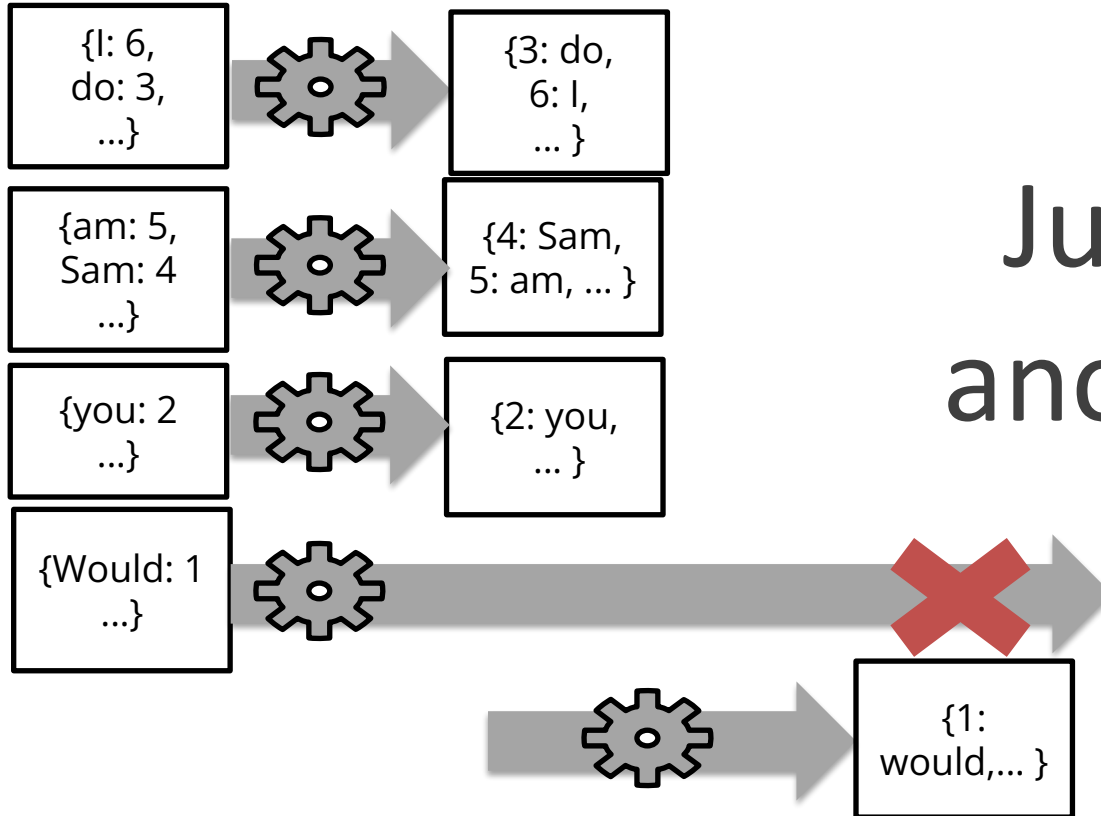
Applied Machine Learning Days '19 [\[link\]](#)

# How to deal with failures?



Just launch another task!

# How to deal with slow tasks?



Just launch  
another task!


# Solution: MapReduce



Jeff Dean

- Smart systems engineers have done all the work for you
  - Task scheduling
  - Virtualization of file system
  - Fault tolerance (incl. data replication)
  - Job monitoring
  - etc.
- **“All”** you need to do: implement Mapper and Reducer classes

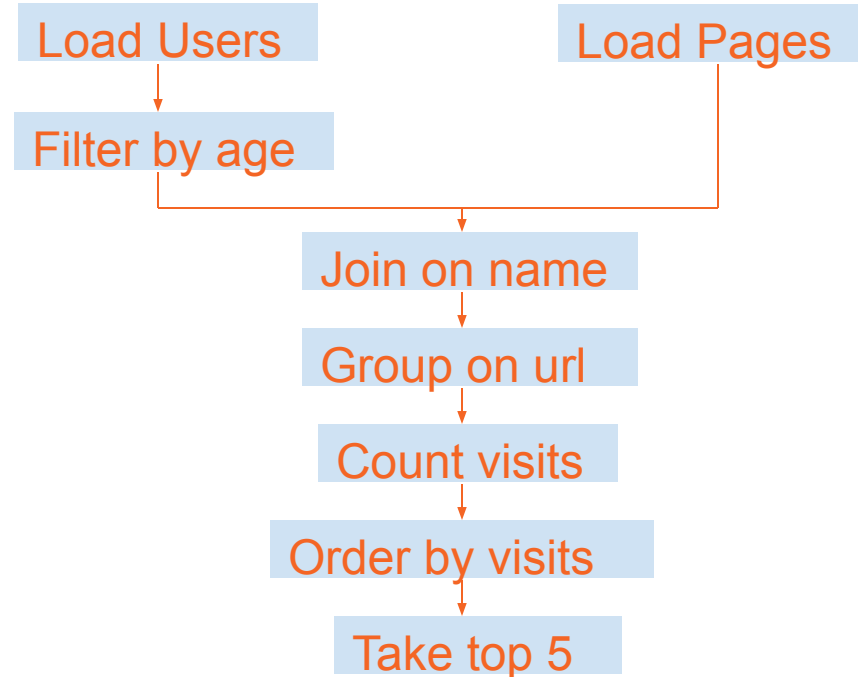
***hadoop***



Need to break more complex jobs into sequence of MapReduce jobs

# Example task

Suppose you have user info in one file, website logs in another, and you need to find the top 5 pages most visited by users aged 18–25



# In MapReduce

[illegible]

# Enter: *Spark*




- A high-level API for programming MapReduce-like jobs

```
sc = SparkContext()
print "I am a regular Python program, using the pyspark lib"
users = sc.textFile('users.tsv') # user <TAB> age
    .map(lambda s: tuple(s.split('\t')))
    .filter(lambda (user, age): age>=18 and age<=25)
pages = sc.textFile('pageviews.tsv') # user <TAB> url
    .map(lambda s: tuple(s.split('\t')))
counts = users.join(pages)
    .map(lambda (user, (age, url)): (url, 1))
    .reduceByKey(add)
    .takeOrdered(5)
```

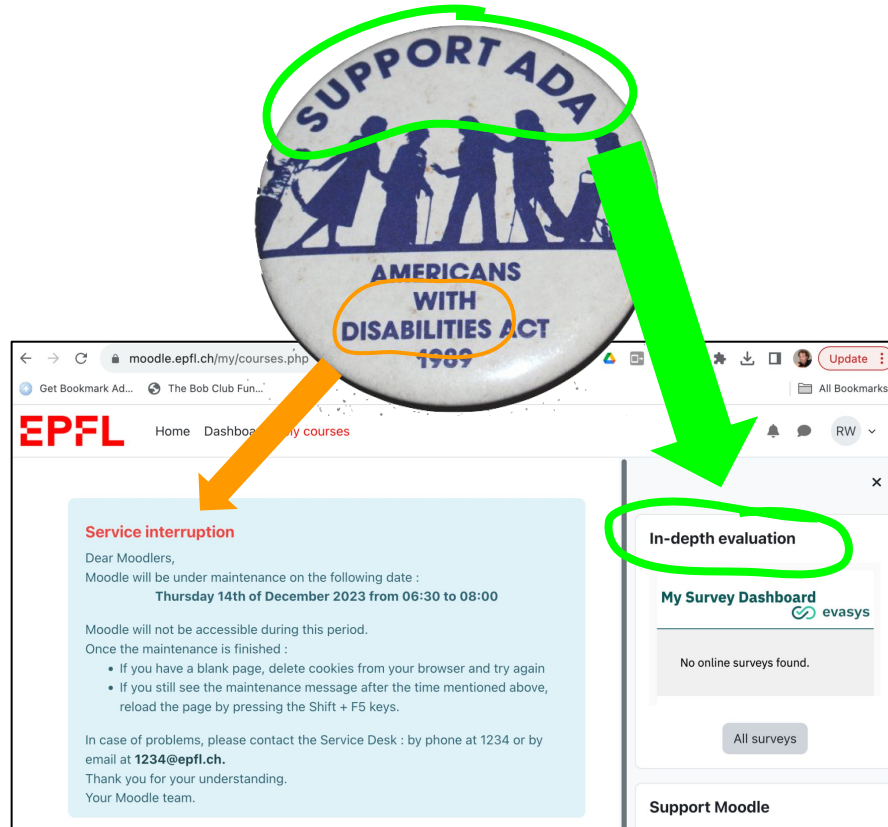




- Implemented in  **Scala** (go EPFL!)
- Additional APIs in
  - Python
  - Java
  - R



# Commercial break



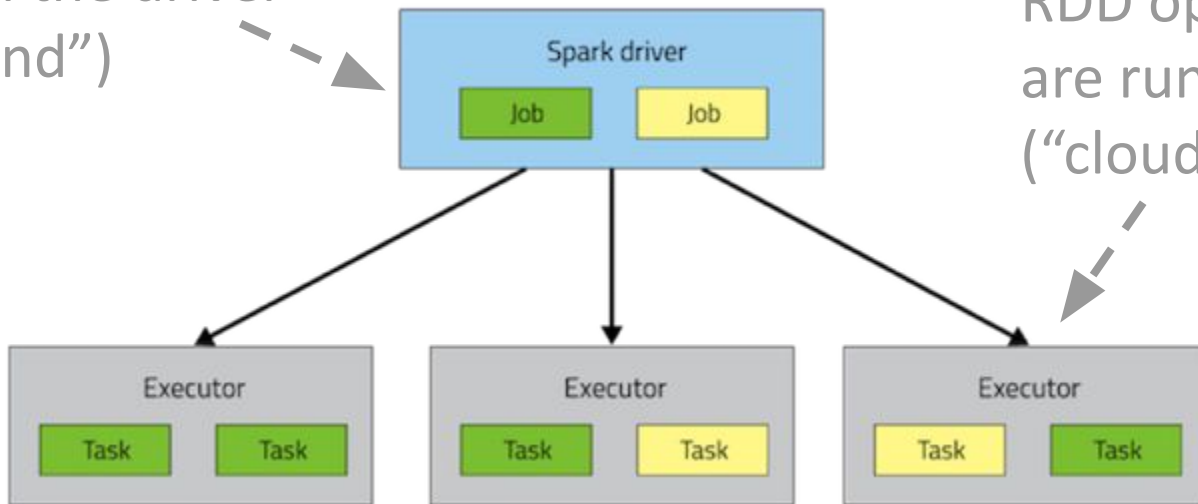
# RDD: resilient distributed dataset



- To programmer: looks like one single list (each element represents a “row” of a dataset)
- Under the hood: oh boy...
  - RDDs “live in the cloud”: split over several machines, replicated, etc.
  - Can be processed in parallel
  - Can be transformed to a single, real list (if small...)
  - Typically read from the distributed file system (HDFS)
  - Can be written to the distributed file system

# Spark architecture

Your Python script runs in the driver ("ground")



RDD operations are run in executors ("cloud")



# RDD operations



- **“Transformations”**

- Input: RDD; output: another RDD
- Everything remains “in the cloud”
- Example: for every entry in the input RDD, count chars
  - `RDD:['I', 'am', 'you'] → RDD:[1, 2, 3]`

- **“Actions”**

- Input: RDD; output: a value that is returned to the driver
- Result is transferred “from cloud to ground”
- Examples: take a sample of entries from RDD and print it on the driver’s shell; or store results to file (local or distributed)

# Lazy execution [unrelated]

- **Transformations** (i.e.,  $\text{RDD} \rightarrow \text{RDD}$  operations) are not executed until it's really necessary (a.k.a. “lazy execution”)
- Execution of transformations triggered by **actions**
- Why?
  - If you never look at the data, there's no point in manipulating it...
  - Smarter query processing possible:  
E.g., `rdd2 = rdd1.map(f1)`  
`rdd3 = rdd2.filter(f2)`  
Can be done in one go — no need to materialize `rdd2`



"I have good news and bad news"

# RDD transformations [\[full list\]](#)

- **map**(*func*): Return a new distributed dataset formed by passing each element of the source through a function *func*
  - $\{1,2,3\}.\text{map}(\text{lambda } x: x*2) \rightarrow \{2,4,6\}$
- **filter**(*func*): Return a new dataset formed by selecting those elements of the source on which *func* returns true
  - $\{1,2,3\}.\text{filter}(\text{lambda } x: x \leq 2) \rightarrow \{1,2\}$
- **flatMap**(*func*): Similar to map, but each input item can be mapped to 0 or more output items (so *func* should return a list rather than a single item)
  - $\{1,2,3\}.\text{flatMap}(\text{lambda } x: [x,x*10]) \rightarrow \{1,10,2,20,3,30\}$

# RDD transformations [\[full list\]](#)

- **sample**(*withReplacement?*, *fraction*, *seed*): Sample a fraction *fraction* of the data, with or without replacement, using a given random number generator *seed*
- **union**(*otherDataset*): Return a new dataset that contains the union of the elements in the source dataset and the argument.
- **intersection**(*otherDataset*): ...
- **distinct**(): Return a new dataset that contains the distinct elements of the source dataset.



# RDD transformations [\[full list\]](#)

- **sample**(*withReplacement?*, *fraction*, *seed*): Sample a fraction *fraction* of the data, with or without replacement, using a given random number generator *seed*

Why *relative fraction*,  
and not *absolute  
number*?

## POLLING TIME

Scan QR code or go to  
<https://web.speakup.info/room/join/66626>



# RDD transformations [\[full list\]](#)

- **groupByKey()**: When called on a dataset of (K, V) pairs, returns a dataset of (K, Iterable<V>) pairs.
  - $\{(1,a), (2,b), (1,c)\}.groupByKey() \rightarrow \{(1,[a,c]), (2,[b])\}$
- **reduceByKey(func)**: When called on a dataset of (K, V) pairs, returns a dataset of (K, V) pairs where the values for each key are aggregated using the given reduce function *func*, which must be of type (V, V) => V.
  - $\{(1, 3.1), (2, 2.1), (1, 1.3)\}.reduceByKey(\text{lambda } (x,y): x+y) \rightarrow \{(1, 4.4), (2, 2.1)\}$

# RDD transformations [\[full list\]](#)

- **sortByKey()**: When called on a dataset of (K, V) pairs, returns a dataset of (K, V) pairs sorted by keys
- **join(*otherDataset*)**: When called on datasets of type (K, V) and (K, W), returns a dataset of (K, (V, W)) pairs with all pairs of elements for each key
  - $\{(1,a), (2,b)\}.join(\{(1,A), (1,X)\}) \rightarrow \{(1, (a,A)), (1, (a,X))\}$
- Analogous: **leftOuterJoin**, **rightOuterJoin**, **fullOuterJoin**
- (There are several other RDD transformations, and some of the above have additional arguments; cf. [tutorial](#))

# RDD actions [\[full list\]](#)

- **collect()**: Return all the elements of the dataset as an array at the driver program. This is usually useful after a filter or other operation that returns a sufficiently small subset of the data.
- **count()**: Return the number of elements in the dataset.
- **take(*n*)**: Return an array with the “first” *n* elements of the dataset.
- **saveAsTextFile(*path*)**: Write the elements of the dataset as a text file in a given directory in the local filesystem or HDFS.
- (There are several other RDD actions; cf. [tutorial](#))

# Broadcast variables

- `my_set = set(range(1e80))`  
`rdd2 = rdd1.filter(lambda x: x in my_set)`  
^ This is a bad idea: `my_set` needs to be shipped with every task (one task per data partition, so if `rdd1` is spread over  $N$  partitions, the above will require copying the same object  $N$  times)
- Better:  
`my_set = sc.broadcast(set(range(1e80)))`  
`rdd2 = rdd1.filter(lambda x: x in my_set.value)`  
^ This way, `my_set` is copied to each executor only once and persists across all tasks (one per partition) on the same executor
- Broadcast variables are **read-only**

# Accumulators

- `def f(x): return x*2`  
`rdd2 = rdd1.map(f)`  
^ How can we easily know how many rows there are in rdd1 (without running a costly reduce operation)?
- Side effects via accumulators!  
`counter = sc.accumulator(0)`  
`def f(x): counter.add(1); return x*2`  
`rdd2 = rdd1.map(f)`
- Accumulators are **write-only** (“add-only”) for executors
- Only driver can read the value: `counter.value`

# RDD persistence

```
rdd2 = rdd1.map(f1)
list1 = rdd2.filter(f2).collect()
list2 = rdd2.filter(f3).collect()
```



rdd1.map(f1)  
transformation is  
executed twice

---

```
rdd2 = rdd1.map(f1)
rdd2.persist()
list1 = rdd2.filter(f2).collect()
list2 = rdd2.filter(f3).collect()
```



Result of rdd1.map(f1)  
transformation is cached  
and reused (can choose  
between memory and  
disk for caching)

# Spark DataFrames



- Bridging the gap between your experience with Pandas and the need for distributed computing
  - RDD = list of rows
  - DataFrame = table with rows and typed columns
- Important to understand what RDDs are and what they offer, but today most of the tasks can be accomplished with DataFrames (**higher level of abstraction  $\Rightarrow$  less code**)
- <https://www.databricks.com/spark/getting-started-with-apache-spark/dataframes>



# Spark SQL [\[link\]](#)



```
sc = SparkContext()
```

```
sqlContext = HiveContext(sc)
```

```
df = sqlContext.sql("SELECT * from table1 GROUP BY id")
```



# Spark's Machine Learning Toolkit

MLlib: Algorithms [[more details](#)]

Classification

- Logistic regression, decision trees, random forests

Regression

- Linear (with L1 or L2 regularization)

Unsupervised:

- Alternating least squares
- K-means
- SVD
- Topic modeling (LDA)

Optimizers

- Optimization primitives (SGD, L-BGFS)

# Example:

## Logistic regression with MLLib

```
from pyspark.mllib.classification \
    import LogisticRegressionWithSGD

trainData = sc.textFile("...").map(...)
testData = sc.textFile("...").map(...)
model = LogisticRegressionWithSGD.train(trainData)
predictions = model.predict(testData)
```

# Remarks

- This lecture is not enough to teach you Spark!
- To use it in practice, you'll need to delve into further online material
- Also: Friday's lab session
- You can't learn it without some frustration :(
- Important skill: assess whether you'd benefit from Spark
  - E.g., >1TB: yes, you'll need Spark
  - 20GB: it depends...



# Feedback

Give us feedback on this lecture here:

<https://go.epfl.ch/ada2023-lec13-feedback>

- What did you (not) like about this lecture?
- What was (not) well explained?
- On what would you like more (fewer) details?
- Where is Waldo?
- ...

# Cluster etiquette

- Develop and debug locally
  - Install Spark locally on your personal computer
  - Use a small subset of the data
- When ready, launch your script on the cluster using [spark-submit](#)
- **Never (never!) use the Spark shell (a.k.a. pyspark) -- it's hereby officially forbidden**
- Useful trench report from a dlab member:  
[“What I learned from processing big data with Spark”](#)